

```
In [46]: import Pkg
Pkg.activate(@__DIR__)
Pkg.instantiate()
import FiniteDiff
import ForwardDiff as FD
import Convex as cvx
import ECOS
using LinearAlgebra
using Plots
using Random
using JLD2
using Test
using MeshCat
const mc = MeshCat
using StaticArrays
using Printf
```

Activating environment at `~/Desktop/Work/Courses/OCRL/HW/HW4_S24/Project.toml`

```
In [47]: include(joinpath(@__DIR__, "utils", "ilc_visualizer.jl"))
```

```
Out[47]: update_car_pose! (generic function with 1 method)
```

Q1: Iterative Learning Control (ILC) (40 pts)

In this problem, you will use ILC to generate a control trajectory for a Car as it swerves to avoid a moose, also known as "the moose test" ([wikipedia \(https://en.wikipedia.org/wiki/Moose_test\)](https://en.wikipedia.org/wiki/Moose_test), [video \(https://www.youtube.com/watch?v=TZ2MYFlnpMI\)](https://www.youtube.com/watch?v=TZ2MYFlnpMI)). We will model the dynamics of the car as with a simple nonlinear bicycle model, with the following state and control:

$$x = \begin{bmatrix} p_x \\ p_y \\ \theta \\ \delta \\ v \end{bmatrix}, \quad u = \begin{bmatrix} a \\ \dot{\delta} \end{bmatrix}$$

where p_x and p_y describe the 2d position of the bike, θ is the orientation, δ is the steering angle, and v is the velocity. The controls for the bike are acceleration a , and steering angle rate $\dot{\delta}$.

```

In [48]: function estimated_car_dynamics(model::NamedTuple, x::Vector, u::Vector)::Vector
or
    # nonlinear bicycle model continuous time dynamics
    px, py, θ, δ, v = x
    a, δdot = u

    β = atan(model.lr * δ, model.L)
    s,c = sincos(θ + β)
    ω = v*cos(β)*tan(δ) / model.L

    vx = v*c
    vy = v*s

    xdot = [
        vx,
        vy,
        ω,
        δdot,
        a
    ]

    return xdot
end
function rk4(model::NamedTuple, ode::Function, x::Vector, u::Vector, dt::Real)
::Vector
    k1 = dt * ode(model, x, u)
    k2 = dt * ode(model, x + k1/2, u)
    k3 = dt * ode(model, x + k2/2, u)
    k4 = dt * ode(model, x + k3, u)
    return x + (1/6)*(k1 + 2*k2 + 2*k3 + k4)
end

```

Out[48]: rk4 (generic function with 1 method)

We have computed an optimal trajectory X_{ref} and U_{ref} for a moose test trajectory offline using this `estimated_car_dynamics` function. Unfortunately, this is a highly approximate dynamics model, and when we run U_{ref} on the car, we get a very different trajectory than we expect. This is caused by a significant sim to real gap. Here we will show what happens when we run these controls on the true dynamics:

```

In [49]: function load_car_trajectory()
    # load in trajectory we computed offline
    path = joinpath(@__DIR__, "utils", "init_control_car_ilc.jld2")
    F = jldopen(path)
    Xref = F["X"]
    Uref = F["U"]
    close(F)
    return Xref, Uref
end
function true_car_dynamics(model::NamedTuple, x::Vector, u::Vector)::Vector
    # true car dynamics
    px, py,  $\theta$ ,  $\delta$ , v = x
    a,  $\delta$ dot = u

    # sluggish controls (not in the approximate version)
    a = 0.9*a - 0.1
     $\delta$ dot = 0.9* $\delta$ dot - .1* $\delta$  + .1

     $\beta$  = atan(model.lr *  $\delta$ , model.L)
    s,c = sincos( $\theta$  +  $\beta$ )
     $\omega$  = v*cos( $\beta$ )*tan( $\delta$ ) / model.L

    vx = v*c
    vy = v*s

    xdot = [
        vx,
        vy,
         $\omega$ ,
         $\delta$ dot,
        a
    ]

    return xdot
end

@testset "sim to real gap" begin
    # problem size
    nx = 5
    nu = 2
    dt = 0.1
    tf = 5.0
    t_vec = 0:dt:tf
    N = length(t_vec)
    model = (L = 2.8, lr = 1.6)

    # optimal trajectory computed offline with approximate model
    Xref, Uref = load_car_trajectory()

    # TODO: simulated Uref with the true car dynamics and store the states in
Xsim
    Xsim = [zeros(nx) for i = 1:N]

    # -----testing-----
    @test norm(Xsim[1] - Xref[1]) == 0
end

```

```


    @test norm(Xsim[end] - [3.26801052, 15.0590156, 2.0482790, 0.39056168, 4.
5],Inf) < 1e-4

    # -----plotting/animation-----
    Xm= hcat(Xsim...)
    Xrefm = hcat(Xref...)
    plot(Xrefm[1,:], Xrefm[2,:], ls = :dash, label = "reference",
        xlabel = "x (m)", ylabel = "y (m)", title = "Simulation vs Referenc
e")
    display(plot!(Xm[1,:], Xm[2,:], label = "actual"))

end

```

```
sim to real gap: Test Failed at In[49]:55
  Expression: norm(Xsim[1] - Xref[1]) == 0
  Evaluated: 13.46291201783626 == 0
Stacktrace:
 [1] macro expansion
      @ In[49]:55 [inlined]
 [2] macro expansion
      @ /buildworker/worker/package_linux64/build/usr/share/julia/stdlib/v1.6/Test/src/Test.jl:1151 [inlined]
 [3] top-level scope
      @ In[49]:39
sim to real gap: Test Failed at In[49]:56
  Expression: norm(Xsim[end] - [3.26801052, 15.0590156, 2.048279, 0.39056168, 4.5], Inf) < 0.0001
  Evaluated: 15.0590156 < 0.0001
Stacktrace:
 [1] macro expansion
      @ In[49]:56 [inlined]
 [2] macro expansion
      @ /buildworker/worker/package_linux64/build/usr/share/julia/stdlib/v1.6/Test/src/Test.jl:1151 [inlined]
 [3] top-level scope
      @ In[49]:39
```



Test Summary:		Fail	Total
sim to real gap		2	2

Some tests did not pass: 0 passed, 2 failed, 0 errored, 0 broken.

Stacktrace:

```
[1] finish(ts::Test.DefaultTestSet)
    @ Test /buildworker/worker/package_linux64/build/usr/share/julia/stdlib/v1.6/Test/src/Test.jl:913
[2] macro expansion
    @ /buildworker/worker/package_linux64/build/usr/share/julia/stdlib/v1.6/Test/src/Test.jl:1161 [inlined]
[3] top-level scope
    @ In[49]:39
```

In order to account for this, we are going to use ILC to iteratively correct our control until we converge.

To encourage the trajectory of the bike to follow the reference, the objective value for this problem is the following:

$$J(X, U) = \sum_{i=1}^{N-1} \left[\frac{1}{2} (x_i - x_{ref,i})^T Q (x_i - x_{ref,i}) + \frac{1}{2} (u_i - u_{ref,i})^T R (u_i - u_{ref,i}) \right] + \frac{1}{2} (x_N - x_{ref,N})^T Q (x_N - x_{ref,N})$$

Using ILC as described in [Lecture 18 \(https://github.com/Optimal-Control-16-745/lecture-notebooks/blob/main/Lecture%2018/Lecture%2018.pdf\)](https://github.com/Optimal-Control-16-745/lecture-notebooks/blob/main/Lecture%2018/Lecture%2018.pdf), we are to linearize our approximate dynamics model about X_{ref} and U_{ref} to get the following Jacobians:

$$A_k = \left. \frac{\partial f}{\partial x} \right|_{x_{ref,k}, u_{ref,k}}, \quad B_k = \left. \frac{\partial f}{\partial u} \right|_{x_{ref,k}, u_{ref,k}}$$

where $f(x, u)$ is our **approximate discrete** dynamics model (`estimated_car_dynamics + rk4`). **You will form these Jacobians exactly once, using X_{ref} and U_{ref}** . Here is a summary of the notation:

- X_{ref} (X_{ref}) - Optimal trajectory computed offline with approximate dynamics model.
- U_{ref} (U_{ref}) - Optimal controls computed offline with approximate dynamics model.
- X_{sim} (X_{sim}) - Simulated trajectory with real dynamics model.
- \bar{U} (\bar{U}) - Control we use for simulation with real dynamics model (this is what ILC updates).

In the second step of ILC, we solve the following optimization problem:

$$\begin{aligned} \min_{\Delta x_{1:N}, \Delta u_{1:N-1}} \quad & J(X_{sim} + \Delta X, \bar{U} + \Delta U) \\ \text{st} \quad & \Delta x_1 = 0 \\ & \Delta x_{k+1} = A_k \Delta x_k + B_k \Delta u_k \quad \text{for } k = 1, 2, \dots, N-1 \end{aligned}$$

We are going to initialize our \bar{U} with U_{ref} , then the ILC algorithm will update $\bar{U} = \bar{U} + \Delta U$ at each iteration. It should only take 5-10 iterations to converge down to $\|\Delta U\| < 1 \cdot 10^{-2}$. You do not need to do any sort of linesearch between ILC updates.



```

In [ ]: # feel free to use/not use any of these

# function trajectory_cost(Xsim::Vector{Vector{Float64}}, # simulated states
#                          Ubar::Vector{Vector{Float64}}, # simulated controls
# (ILC iterates this)
#                          Xref::Vector{Vector{Float64}}, # reference X's we want to track
#                          Uref::Vector{Vector{Float64}}, # reference U's we want to track
#                          Q::Matrix,                    # LQR tracking cost term
#                          R::Matrix,                    # LQR tracking cost term
#                          Qf::Matrix                    # LQR tracking cost term
#                          )::Float64                    # return cost J

#     J = 0
#     J += 0.5 * cvx.quadform(Xsim[end] - Xref[end], Qf)
#     # TODO: return trajectory cost J(Xsim, Ubar)
#     for i = 1:length(Xsim)-1
#         J += 0.5*cvx.quadform(Xsim[i] - Xref[i], Q) + 0.5*cvx.quadform(Ubar[i] - Uref[i], R)
#     end
#     return J
# end

function vec_from_mat(Xm::Matrix)::Vector{Vector{Float64}}
    # convert a matrix into a vector of vectors
    X = [Xm[:,i] for i = 1:size(Xm,2)]
    return X
end

function ilc_update(Xsim::Vector{Vector{Float64}}, # simulated states
                   Ubar::Vector{Vector{Float64}}, # simulated controls (ILC iterates this)
                   Xref::Vector{Vector{Float64}}, # reference X's we want to track
                   Uref::Vector{Vector{Float64}}, # reference U's we want to track
                   As::Vector{Matrix{Float64}},   # vector of A jacobians at each time step
                   Bs::Vector{Matrix{Float64}},   # vector of B jacobians at each time step
                   Q::Matrix,                      # LQR tracking cost term
                   R::Matrix,                      # LQR tracking cost term
                   Qf::Matrix                       # LQR tracking cost term
                   )::Vector{Vector{Float64}}      # return vector of ΔU's

    # solve optimization problem for ILC update
    N = length(Xsim)
    nx,nu = size(Bs[1])

    # create variables
    ΔX = cvx.Variable(nx, N)
    ΔU = cvx.Variable(nu, N-1)

```

```

# TODO: cost function (tracking cost on Xref, Uref)
cost = 0.5*cvx.quadform( $\Delta X[:, \text{end}] + X_{\text{sim}}[\text{end}] - X_{\text{ref}}[\text{end}]$ , Qf)
for i = 1:N-1
    cost += 0.5*cvx.quadform( $\Delta X[:, i] + X_{\text{sim}}[i] - X_{\text{ref}}[i]$ , Q) + 0.5*cvx.qua
dform( $\Delta U[:, i] + U_{\text{bar}}[i] - U_{\text{ref}}[i]$ , R)
end

# problem instance
prob = cvx.minimize(cost)

# TODO: initial condition constraint
prob.constraints += ( $\Delta X[:, 1] == \text{zeros}(nx)$ )
# TODO: dynamics constraints
for i = 1:N-1
    prob.constraints += ( $\Delta X[:, i+1] == A_s[i]*\Delta X[:, i] + B_s[i]*\Delta U[:, i]$ )
end

cvx.solve!(prob, ECOS.Optimizer; silent_solver = true)

# return  $\Delta U$ 
 $\Delta U = \text{vec\_from\_mat}(\Delta U.\text{value})$ 

return  $\Delta U$ 
end

```

Out[]: ilc_update (generic function with 1 method)

Here you will run your ILC algorithm. The resulting plots should show the simulated trajectory X_{sim} tracks X_{ref} very closely, but there should be a significant difference between U_{ref} and U_{bar} .

In [50]: @testset "ILC" begin

```
# problem size
nx = 5
nu = 2
dt = 0.1
tf = 5.0
t_vec = 0:dt:tf
N = length(t_vec)

# optimal trajectory computed offline with approximate model
Xref, Uref = load_car_trajectory()

# initial and terminal conditions
xic = Xref[1]
xg = Xref[N]

# LQR tracking cost to be used in ILC
Q = diagm([1,1,.1,.1,.1])
R = .1*diagm(ones(nu))
Qf = 1*diagm(ones(nx))

# load all useful things into params
model = (L = 2.8, lr = 1.6)

params = (Q = Q, R = R, Qf = Qf, xic = xic, xg = xg, Xref=Xref, Uref=Uref,
          dt = dt,
          N = N,
          model = model)

# this holds the sim trajectory (with real dynamics)
Xsim = [zeros(nx) for i = 1:N]

# this is the feedforward control ILC is updating
Ubar = [zeros(nu) for i = 1:(N-1)]
Ubar .= Uref # initialize Ubar with Uref

# TODO: calculate Jacobians
A = [FD.jacobian(x -> rk4(model, estimated_car_dynamics, x, Uref[i], dt),
Xref[i]) for i = 1:N-1]
B = [FD.jacobian(u -> rk4(model, estimated_car_dynamics, Xref[i], u, dt),
Uref[i]) for i = 1:N-1]

# logging stuff
@printf "iter      objv      |ΔU|      \n"
@printf "-----\n"

for ilc_iter = 1:10 # it should not take more than 10 iterations to converge
    Xsim[1] = xic
    # TODO: rollout
    for i = 1:N-1
        Xsim[i+1] = rk4(model, true_car_dynamics, Xsim[i], Ubar[i], dt)
    end
    # TODO: calculate objective val (trajectory_cost)
```

```

    obj_val = 0.5*Xsim[end]*Qf*Xsim[end]
    for i = 1:N-1
        obj_val += 0.5*(Xsim[i] - Xref[i])*Q*(Xsim[i] - Xref[i]) + 0.5*(Ubar[i] - Uref[i])*R*(Ubar[i] - Uref[i])
    end

    # solve optimization problem for update (ilc_update)
    ΔU = ilc_update(Xsim, Ubar, Xref, Uref, A, B, Q, R, Qf)

    # TODO: update the control
    Ubar = Ubar .+ ΔU

    # logging
    @printf("%3d    %10.3e    %10.3e    \n", ilc_iter, obj_val, sum(norm.(ΔU)))

end

# -----plotting/animation-----
Xm= hcat(Xsim...)
Um = hcat(Ubar...)
Xrefm = hcat(Xref...)
Urefm = hcat(Uref...)
plot(Xrefm[1,:], Xrefm[2,:], ls = :dash, label = "reference",
      xlabel = "x (m)", ylabel = "y (m)", title = "Trajectory")
display(plot!(Xm[1,:], Xm[2,:], label = "actual"))

plot(t_vec[1:end-1], Urefm', ls = :dash, lc = [:green :blue], label = "",
      xlabel = "time (s)", ylabel = "controls", title = "Controls (-- is reference)")
display(plot!(t_vec[1:end-1], Um', label = ["δ" "a"], lc = [:green :blue]))

# animation
vis = Visualizer()
vis_traj!(vis, :traj, [[x[1],x[2],0.1] for x in Xsim]; R = 0.02)
build_car!(vis[:car])
anim = mc.Animation(floor(Int,1/dt))
for k = 1:N
    mc.atframe(anim, k) do
        update_car_pose!(vis[:car], Xsim[k])
    end
end
mc.setanimation!(vis, anim)
display(render(vis))


# -----testing-----
@test 0.1 <= sum(norm.(Xsim - Xref)) <= 1.0 # should be ~0.7
@test 5 <= sum(norm.(Ubar - Uref)) <= 10 # should be ~7.7

end


```

iter	objv	$ \Delta U $
1	1.409e+03	6.307e+01
2	1.193e+03	4.498e+01
3	7.406e+02	9.266e+01
4	1.082e+02	1.394e+01
5	9.529e+01	1.959e+00
6	9.094e+01	1.679e-01
7	9.050e+01	1.649e-02
8	9.045e+01	1.578e-03
9	9.044e+01	1.911e-04
10	9.044e+01	3.029e-05

Trajectory



Controls (-- is reference)



[Info: Listening on: 127.0.0.1:8700, thread id: 1

└ Info: MeshCat server started. You can open the visualizer by visiting the following URL in your browser:

└ <http://127.0.0.1:8700>

4

Test Summary:	Pass	Total
ILC	2	2

Out[50]: Test.DefaultTestSet("ILC", Any[], 2, false, false)